

Experiment No : 04

Date of Experiment : 24 / 04 / 2025

Experiment Name : Round Robin Scheduling Algorithm

Theory:

Round Robin (RR) is a **preemptive** CPU scheduling algorithm that assigns a **fixed time slice** or **time quantum** to each process in the ready queue. If a process does not finish execution within that quantum, it is **preempted** and placed at the **end of the queue**. This continues cyclically until all processes are completed. This scheduling method is **fair and efficient**, especially in **time-sharing systems**, as it prevents any single process from **monopolizing the CPU**. However, the performance of the algorithm heavily depends on the **size of the time quantum**. If the quantum is too small, the system experiences excessive **context switching**. If too large, it behaves similarly to FCFS.

Code:

```
print("ROUND ROBIN SCHEDULING")
n = int(input("Enter number of processes: "))
processes = []
for i in range(n):
    pid = "P" + str(i + 1)
    arrival = int(input(f"Enter arrival time of process {i+1}: "))
    burst = int(input(f"Enter burst time of process {i+1}: "))
    processes.append([pid, arrival, burst])
time_quantum = int(input("Enter Time Quantum: "))
current_time = 0
queue = []
completed = []
remaining_burst = {}
arrival_map = {}
for p in processes:
    pid, arrival, burst = p
    remaining_burst[pid] = burst
    arrival_map[pid] = arrival
processes.sort(key=lambda x: x[1])
ready = processes.copy()
visited = set()
while True:
    for p in ready:
        if p[1] <= current_time and p[0] not in visited:
            queue.append(p[0])
            visited.add(p[0])
    if queue:
        pid = queue.pop(0)
        arrival = arrival_map[pid]
        burst = remaining_burst[pid]
        exec_time = min(time_quantum, burst)
        current_time += exec_time
```

```

remaining_burst[pid] -= exec_time
for p in ready:
    if p[1] <= current_time and p[0] not in visited:
        queue.append(p[0])
        visited.add(p[0])
    if remaining_burst[pid] > 0:
        queue.append(pid)
    else:
        for p in processes:
            if p[0] == pid:
                finish = current_time
                tat = finish - p[1]
                wt = tat - p[2]
                completed.append([pid, p[1], p[2], finish, tat, wt])
        else:
            if len(completed) == n:
                break
            current_time += 1
avg_wt = sum(p[5] for p in completed) / n
print("\nProcess | Arrival | Burst | Exit | Turn Around | Wait |")
for p in completed:
    print(f' {p[0]} | {p[1]} | {p[2]} | {p[3]} | {p[4]} | {p[5]} ')
print(f'\nAverage Waiting Time: {avg_wt:.2f}')

```

Input and Output:

```

Run osLab4 x
D:\SQL\.venv\Scripts\python.exe "D:\0parating system\osLab4.py"
ROUND ROBIN SCHEDULING
Enter number of processes: 4
Enter arrival time of process 1: 0
Enter burst time of process 1: 5
Enter arrival time of process 2: 1
Enter burst time of process 2: 3
Enter arrival time of process 3: 2
Enter burst time of process 3: 1
Enter arrival time of process 4: 4
Enter burst time of process 4: 2
Enter Time Quantum: 2

Process | Arrival | Burst | Exit | Turn Around | Wait |
P3      | 2       | 1     | 5    | 3           | 2
P4      | 4       | 2     | 9    | 5           | 3
P2      | 1       | 3     | 10   | 9           | 6
P1      | 0       | 5     | 11   | 11          | 6

Average Waiting Time: 4.25

```

Explanation:

- **Input Handling:** The user enters the number of processes, their arrival time, and burst time, followed by the time quantum.
- **Ready Queue Maintenance:** Processes are added to the ready queue as they arrive.
- **Execution Logic:** Each process is executed for a maximum of one time quantum. If unfinished, it is placed back in the queue.
- **Completion Time:** Updated after each time slice execution.
- **Turnaround Time (TAT):** Calculated as Completion Time – Arrival Time.
- **Waiting Time (WT):** Calculated as Turnaround Time – Burst Time.
- **Output:** A summary table and average waiting time are printed at the end.

Discussion:

The Round Robin algorithm fairly allocates CPU time to each process and **prevents starvation**. It is ideal for systems requiring **frequent user interaction**. In this implementation:

1. **Arrival time** is considered, making it realistic for modern OS behavior.
2. The **time quantum** influences performance:
 - A small quantum leads to **high responsiveness** but **more context switching**.
 - A large quantum may behave like FCFS.

Conclusion:

Round Robin Scheduling is a fair and straightforward algorithm suited for **interactive and time-sharing** systems. It ensures **equal CPU access** and avoids starvation. However, the **time quantum must be carefully chosen** to maintain an efficient balance between context switching and process throughput. It remains a foundational scheduling strategy in modern operating systems.